

TagFS - A Tag Based Filesystem

Catalina Macalet

Computer Science Department
Politehnica University of Bucharest
Email: catalina.macalet@cti.pub.ro

Eugen Hristev

Computer Science Department
Politehnica University of Bucharest
Email: eugen.hristev@cti.pub.ro

Mihai Dinu

Computer Science Department
Politehnica University of Bucharest
Email: mihai.dinu@cti.pub.ro

Sorin Dumitru

Computer Science Department
Politehnica University of Bucharest
Email: sorin.dumitru@cti.pub.ro

Abstract—File systems are an integral part of every operating system. Because of the high capacity of modern hard drives file systems need a better way of organizing and accessing data.

TagFS implements a tag-based filesystem in Linux which offers support for tagging files and browsing files by tags.

Index Terms—Tags, file systems, VFS

I. INTRODUCTION

In most operating systems the files are hierarchically organized. This means that there usually is a starting point, or parent directory. In Windows based systems there are multiple starting points based on the physical hard drive partitions. In Unix-like systems, there is a single root drive with different mount points available for users to add or remove subtrees from different drives, partitions, etc. In these filesystems a user organizes related data by storing it in the same folder but say that a user, Bob, has two separated folders one for storing photos taken in the mountains (Mountain-pics) and one for storing photos in which a certain person appears (Alice-pics). Two questions arise, one, where should Bob store a picture taken in the mountains in which Alice appears and two, how could Bob find the pictures taken in the mountains in which Alice appears. For current filesystems the answer to the first question might be storing the photo in either folder and in the second one creating a link to this photo or, store it in both folders. The answer to the second one could be naming the photo in such a way that retrieving them based on the previously stated criteria would work. TagFS file system aims to bring a different approach, based on tags rather than hierarchical system that is rooted for a long time in modern operating system. For the above example, for a TagFS filesystem the answer to both questions would be adding tags to photos (<mountains><Alice>) and then search for files that contain these tags. The question of where to store a specific photo would not be that important anymore. A pure tag file system is difficult to implement starting from zero, so we tried to adapt the current file system in Linux to support tags and see how the two systems can coexist on an end-user machine. A tag file system should be able to organize files, data on the disk regardless of hierarchical logical approach. The position of the files on the disk is

irrelevant and completely transparent to the user. The file system should be able to put files on disks and simply recover them on demand based on tags requests. In our approach, logical directory based organization and file tags coexist, in order to see how the two systems can fit and how the user can use alternatives for searching and clustering the information it has. We implemented a tag layer in the Linux Virtual File System and tested how this impacts the regular user, we added possibilities for the user to play with tags: add, delete, search.

II. STATE OF THE ART

The idea of tagging files in order to access them in an easier fashion is not a new one and various attempts to implement solutions have been made. Some of these are specialized solutions for special kind of data, such as Calibre which makes ebook management easier, implemented in userspace. The vast majority of these applications rely on a database where mappings between files and associated metadata are stored and expose a set of commands which translate to specific queries for the database.

A. Nepomuk-KDE

¹ Nepomuk-KDE is an implementation of Nepomuk which has been integrated with KDE and that allows adding metadata to items stored on a computer and making queries based on that metadata. Dolphin KDE file manager allows adding and removing basic metadata to files, such as comments and tags.

TODO:

more on this

B. Other Projects

TagFS: Bringing Semantic Metadata to the Filesystem² is a research project started at the University of Koblenz which, as Nepomuk, relies on RDF for defining semantics and SPARQL. Metadata is stored in a repository having an associated graph, and various operations can be performed on it (additions, updates, etc).

TODO:

existing implementations for tag based filesystems.

¹ <http://nepomuk.kde.org/>

² <http://www.eswc2006.org/poster-papers/FP31-Schenk.pdf/>

TODO:

Add details about how tags are stored and browsed

B. Storage

TODO:

RAM storage, picture, details.

C. Tag-handling commands

TODO:

pic with command flow

1) *Implementation details:*

- *VFS Hooks*
loreipsumloreipsumloreipsumloreipsumloreipsumloreipsumloreipsum
- *Userspace application*
loreipsumloreipsumloreipsumloreipsumloreipsumloreipsumloreipsum

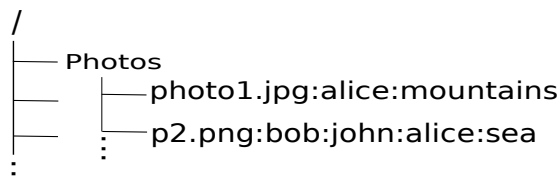


Fig. 1. TagFS hierarchy

The TagFS application architecture is presented in Figure 2. The *CLI* is used to issue commands for tag manipulation.




Fig. 2. TagFS architecture

There are two types of commands. The first type consists of file manipulation commands available on every Unix-like operating system, such as *ls*, *touch*, *mv*, *cp* whose behaviour and implementation was changed in TagFS implementation in order to support tags. The second type of commands refers to new TagFS commands implemented in order to allow more tag-related operations. The implementation changes are related to hooks created in *VFS* and will be detailed in subsection *Tag handling*. No implementation changes at filesystem level were required.

A. Architectural decision

TagFS started as an idea to create a more user-friendly file system; remembering tags is easier than remembering the name of a file or the place where it is stored but at the same time a pure tag filesystem would offer no simple way of organizing data in a hierarchical manner, a choice to implement TagFS as a new filesystem, from scratch, thus would have been pointless. The other choice was to implement TagFS as hooks in VFS in order to store and retrieve needed metadata. Since changes are made at VFS level there will be an overhead for filesystems that subsequently are to be used without tag support. A main concern in implementation was to reduce this overhead to as little as possible.

IV. CONCLUSION

The conclusion goes here.

Possible future work: In a pure tag file system, the disk mechanism could be improved in the following way: We know that tags can be added to some files, we have no hierarchical structure of the files. This way we can find blocks of files based on tags which could reduce disk fragmentation. Clustering tag data can give insight on how much space there is required of a certain tag type files and how accessible this should be to the user. This could lower the external fragmentation of the disk if properly used. However more tests should be done regarding this problem.

-graphical if

REFERENCES

- [1] Stephan Bloehdorn and Max Volkel, Tagfs-tag semantics for hierarchical file systems,
- [2] Yuan-Liang Tai1, Shang-Rong Tsai, Guang-Hung Huang, Chia-Ming Lee, Lian-Jou Tsai, Kuo-Feng Ssu and Shou-Jen Wey, A Label-Based File System
- [3] Nepomuk-KDE <http://nepomuk.kde.org/>